

Project Sanctum - Korvin Merrow World Construction Transcript

Project: Project Sanctum: Clinical Evaluation Benchmark
World Name: Korvin Merrow (formerly James Carter)
Patient: 62-year-old male, complex comorbidity presentation
Supporting Provenance: <https://claude.ai/share/d5129364-5d6c-4a2c-acb3-282f367a0040>

Table of Contents

- Phase 1: Brainstorm & Alignment
 - Project Readiness Audit
 - Brainstorm Architecture Interview
 - Brainstorm Reviewer Simulation
 - Brainstorm Approval
- Phase 2: World Specification
 - World Spec Architecture
 - World File Ecosystem (22 files)
 - Task-Context Layers (7 files)
 - Supplementary References (4 files)
- Phase 3: Task Layer Construction
 - Prompts (TP-KM01-06)
 - Expected Outputs (EO-KM01-06)
 - Golden Responses (Golden-KM01-06)
 - Grader Guidance (GG-KM01-06)
- Phase 4: Validation & Lock
 - Golden Architecture Validation
 - Grader Guidance Construction Validation
 - Audit Findings & Reconciliations
 - Lock-Readiness Confirmation
 - Forward Dependencies
 - AutoQC Requirements
 - Packaging Carry-Forwards

Phase 1: Brainstorm & Alignment

1.1 Project Readiness Audit

Date: May 29, 2026

Objective: Validate physician-Claude alignment before brainstorm construction

Scope: 8-dimension alignment check across project understanding, role boundaries, workflow state, brainstorm structure, core concepts, AI usage boundaries, active world review, and quality gates

Project Understanding

Project Sanctum evaluates physician-level clinical judgment in frontier AI, not textbook recall. The unit of evaluation is a complete clinical environment (a *World*) with messy documentation, an evolving timeline, competing perspectives, and imperfect information.

Strong World characteristics: - Common conditions interacting through messy charts - Multiple documents that must be reconciled - Competing consultants and perspectives - Source hierarchy matters - Medication changes carry consequences - Discharge safety stays uncertain despite objective improvement - Solvable by a seasoned clinician from provided files alone - Hard for AI that fails to synthesize or prioritize

Weak World characteristics (to avoid): - Rare-diagnosis stack - One-document answers - Pure factual recall - Prompts that leak the trick - Generic complexity with no chart mechanism

Role Boundaries

Dr. Udeogaranya (Physician Expert): - Sole source of clinical truth - Scenario concept ownership - Trap identification - Task idea generation - Clinical/diagnostic reasoning - Medication decisions and management plans - Task prompts, golden responses, and grader guidelines (100% physician-authored)

Claude (Assistant): - Drafting and structuring support - Consistency checking - Formatting and organization - Missing-detail flagging - Reviewer-risk auditing - Critique against Sanctum checklists - QC of human-written prompts/goldens/graders (verification only)

Current Workflow State

Workflow Phase: Phase 1 - World Building onboarding (Steps 1-6)

Pass: Pass 1 - Brainstorm clinical interview

Completed: World Setup interview (substantially complete)

Next: Frictions interview â†’ Traps interview â†’ Rough Task Ideas interview

Required Brainstorm Structure

World Setup: Patient and scenario coherence, setting clarity, specialty/encounter type definition, timeline shape, realistic comorbidity stack, complexity from workflow (not rare diagnosis)

Major Friction Points: True people/perspective conflicts (not mere information gaps), named parties, defensible positions on both sides, clinical significance

Major Traps: Information problems in files (copy-forward errors, buried findings, source-of-truth conflicts, temporal mismatches), trap type identification, mechanism clarity, synthesizable across 2+ documents

Rough Task Ideas: 5-8 concrete independent tasks, deliverable types, workflow mapping, world elements, trap/friction linkage, temporal anchoring post-snapshot

Core Concepts Defined

Scenario: Patient story and clinical journey

World: Complete clinical environment (chart ecosystem, documents, timeline, competing perspectives, source conflicts, information noise)

Task: Realistic clinician workflow using world files; independent encounters not dependent on other tasks

Friction: People/perspective conflict requiring a named party advocating a position

Trap: Information problem in files requiring synthesis across 2+ documents

Discipline: Documentation conflict = trap UNLESS a person/role actively advocates a position due to it â†’ then friction

AI Usage Boundaries

Claude may help with: - Brainstorming support built on physician material - Boilerplate and structured drafting - Consistency checking and formatting - Missing-detail identification - Reviewer-risk auditing - Critique against Sanctum checklists

Claude must not independently create: - Scenario concept - Any trap - Any task idea - Clinical/medical reasoning - Task prompts - Golden responses - Grader guidelines (100% physician-authored; QC only)

Active World Review: Korvin Merrow Concept

Patient Profile:

62-year-old male with long-standing type 2 diabetes, hypertension, CKD stage 3, HFrEF, CAD, polypharmacy, and PMR history on chronic prednisone with recent taper

Presentation:

ED arrival with altered mental status, progressive weakness, poor oral intake, borderline hypotension, AKI on CKD, family-reported urinary symptoms/subjective fever

Working Diagnosis & Course:

Initial: urinary-source sepsis (appropriate treatment)

Course evolution: infection improves over multi-day admission; not all symptoms resolve

Contributors requiring reassessment: infection, dehydration/AKI, medication effects, endocrine/steroid issues, chronic disease progression

Snapshot Timing:

Hospital day 5-7 in discharge planning

Vitals improved, infection controlled, renal function improving

Patient still weaker than baseline with intermittent confusion

Significantly changed medication plan

Core Question:

Medically stable on paper vs. actually safe to leave?

Locked Clinical Decisions

Patient profile and comorbidities

Presentation and initial urosepsis working diagnosis

Initial sepsis treatment appropriateness

Course shape (infection improves, residual symptoms remain, multiple contributors reassessed)

Steroid thread as reconciliation problem (not hidden reveal)

Baseline function (independent ADLs, ambulatory, family-supported)

Discharge state and day 5-7 snapshot

Unresolved Physician Decisions

Exact frictions (ED/inpatient vs endocrinology, nephrology vs cardiology medication, family discharge requests)

Final trap list with source locations and synthesis chains

Final 5-8 task ideas with workflow mapping and P0 identification

Possible Reviewer Risks

Adrenal insufficiency reading as a hidden zebra

Cardiology involvement weak unless HFREF/CAD medication tension is concrete

Frictions collapsing into mere information gaps

Traps solvable from single document

Tasks too generic or not approved workflows

Vague discharge-safety concern without explicit baseline/residual deficits

Quality Gate: Textbook vs. Complexity

Textbook Diagnosis Puzzle Risk: LOW (contingent)

- Protection: Multi-causal discharge picture (weakness/confusion from deconditioning, residual infection, AKI/metabolic derangement, medication effects, relative adrenal insufficiency from taper, or chronic-disease progression) - Must remain judgment-based; no single document names the answer - Tasks ask for work products (notes, med rec, discharge summary) not “diagnosis” - Failure mode to avoid: steroid timeline as clean breadcrumb trail â†’ must keep scattered and contradictory

Rare Disease Reveal Risk: LOW

- All conditions common in older comorbid adult (PMR-on-steroids, steroid-related adrenal suppression)

- Only becomes zebra if adrenal piece framed as secret solution - Standard already forbids this (“one competing concern, not the whole answer”)

Too Easy for AI Risk: MODERATE (watch-item)

- Agents excel at reading one authoritative summary - Conceptually protected (no single-file traps,

enforced source hierarchy, cross-document synthesis required) - Real defense built in next interviews: source-hierarchy conflict, steroid timeline forcing temporal reasoning, traps needing 2 documents, prompts that don't leak

Readiness Gate: READY - proceed to Frictions interview

1.2 Brainstorm Architecture Interview

Date: May 29, 2026

Objective: Physician-authored brainstorm draft from interview decisions

Frictions Finalized

F1 - Nephrology vs Cardiology

Stakeholders: Both specialties

Nephrology position: Renal/hemodynamic safety prioritized in medication tapering

Cardiology position: Long-term CV protection (beta-blocker/ACE-I continuation)

Assessment: PASS - genuine perspective conflict, both defensible positions

F2 - Family vs Inpatient Medicine

Stakeholders: Patient's wife (Mara Merrow) vs hospitalist team

Family position: Functional readiness and home safety required before discharge

Inpatient team position: Objective physiologic stabilization as discharge marker

Assessment: PASS - valid friction type, concrete positions

F3 - ED/Inpatient Medicine vs Endocrinology (CONDITIONAL)

Stakeholders: Inpatient team vs endocrinology

Inpatient position: Avoid over-continuation of prednisone

Endocrinology position: Adrenal evaluation and taper-safety monitoring required

Assessment: CONDITIONAL PASS - survives as friction only if endocrinology actively advocates in world (recommendation artifact); becomes trap if documentation merely ambiguous

1.3 Brainstorm Reviewer Simulation

Date: May 29, 2026

Reviewer Role: Strict Sanctum standards applied

Scope: World Setup, Frictions, Traps, Task Ideas, Failure Prediction, Approval Verdict

World Setup Verdict: STRONG

Setting clarity: Yes - EM/IM acute hospital, ED + inpatient + consultants + discharge planning, multi-day course, snapshot fixed at day 5-7

Patient realism: Yes - comorbidity stack (T2DM, HTN, CKD3, HFrEF/CAD, polypharmacy, PMR-on-prednisone) internally consistent; geriatric/IM admission archetype; presentation (AMS, weakness, poor intake, possible UTI, AKI) bread-and-butter; baseline (independent ADLs, ambulatory, family-supported) coherent and load-bearing

Complexity source: Yes - common conditions interacting through messy documentation, evolving course, competing consultants, ambiguous discharge; NOT rare diagnosis

Textbook risk: Low at concept level; steroid/adrenal thread could read as zebra but protection asserted (not yet instantiated); flagged for spec-phase instantiation

Verdict: Strong setup, no send-back

Frictions Verdict: PASS (with condition on F3)

F1 - Nephrology vs Cardiology: PASS

Both stakeholders named, both positions reasonable, genuine perspective conflict, strongest of three

F2 - Family vs Inpatient Medicine: PASS

Both sides defensible, valid friction type, concrete positions

F3 - ED/Inpatient vs Endocrinology: CONDITIONAL PASS

Risk: collapsing into trap if endocrinology only provides documentation ambiguity (not advocacy). Must be confirmed in spec as actual recommendation artifact.

Traps Verdict: STRONG

5 identified traps, all clinically plausible, all discoverable from files, all cause meaningful error if missed, all require multi-document synthesis:

Prednisone Source-of-Truth (world-level) - Old rheumatology notes vs. newer outpatient plan vs. med-rec confusion

HF/AKI Medication Reconciliation (world-level) - Cardiology vs nephrology tension unresolved

Buried Functional/Cognitive Evidence (world-level) - Functional status scattered across nursing, PT, OT, family reports

Sepsis Anchoring After Partial Improvement (world-level) - Infection improves but not all symptoms â†’ bias toward infection as explanation

Discharge Source-Hierarchy (task-level) - FI-W22 visible-but-incomplete vs. FI-W series depth

Rough Task Ideas: 5-8 Concrete Tasks

All tasks independent, workflow-based, deliverable-specific, trap/friction-linked, temporally anchored post-snapshot

Approval Verdict: PASS

1.4 Brainstorm Approval

Reviewer Decision: APPROVED by Stacey S (human reviewer)

Feedback: Alignment confirmed, readiness gate passed

Proceed: World Spec phase

Phase 2: World Specification

2.1 World Spec Architecture

Date: May 30, 2026 onwards

Scope: Complete world specification including identity, conditions, medications, baseline anchors, timeline, provider roster, 22 world-level files, 7 task-context files, 4 supplementary files

World Identity

Patient: Korvin Merrow, 62-year-old male

Setting: Acute hospital, EM/IM

Encounter: ED â†’ inpatient admission â†’ consultants â†’ discharge planning

Timeline: 5-7 hospital days, snapshot frozen at day 5-7

Core Problem: Altered mental status, weakness, poor intake, AKI on CKD, sepsis work-up

14 Baseline Conditions

Type 2 Diabetes Mellitus

Hypertension

Chronic Kidney Disease Stage 3 (baseline Cr ~2.0)

Heart Failure with Reduced Ejection Fraction (EF 25-30%)

Coronary Artery Disease (s/p remote PCI/stent)

Polymyalgia Rheumatica (on chronic prednisone)

Obstructive Sleep Apnea

Osteoporosis

Anemia of chronic kidney disease

Hyperparathyroidism (secondary to CKD)

Dyslipidemia

GERD

Benign Prostatic Hyperplasia

Orthostatic Hypotension

20 Baseline Medications (Outpatient)

Detailed list with doses, source authority, and timing (pre-admission state for reconciliation purposes)

Temporal Architecture

Hour 0: ED arrival (acute decompensation)

HD 1-2: ED evaluation, admission, initial workup, sepsis protocol

HD 3: Infection improving trend noted, reassessment of contributors

HD 4: Continued improvement, medication adjustments, consultant input

HD 5-7: Discharge planning, readiness assessment, safety coordination

Snapshot: Fixed at 05/23 18:00 (HD 5-7 timeline, pre-discharge, no post-close leakage)

Source Hierarchies

Authority Hierarchy (for clinical decision-making):

Hospitalist > Consultants > PT/OT > CM/SW > Family > Patient

Master Medication Hierarchy (for med reconciliation):

Attending > Med-Rec > Pharmacy > Consultants > PCP > Family > Patient

Prednisone-Specific Hierarchy (for steroid source):

Rheumatology > Med-Rec > Pharmacy > Family > Patient

2.2 World-Level Files (FI-W01-W22)

22 files total, all locked and evidence-bearing

File

Title

Timeline

Key Content

FI-W01

ED Triage Initial Intake

Hour 0

Chief complaint, vital signs, presentation

FI-W02
ED Provider Assessment
Hour 0-2
Initial H&P, differential, sepsis workup
FI-W03
Admission H&P
HD 1
Comprehensive baseline, 14 conditions, 20 meds, exam
FI-W04
Initial Medication Reconciliation Note
HD 1
Med-rec form, reconciliation conflicts begin
FI-W05
Pharmacy Refill History Report
HD 1
Outpatient fill records, adherence hints, source conflict
FI-W06
Outpatient Rheumatology Prednisone Provenance
Pre-admission
Highest-authority prednisone source
FI-W07
Primary Care Outpatient Baseline Summary
Pre-admission
PCP perspective, medication list version
FI-W08-W11
Hospitalist Progress Notes
HD 1-6
Daily course documentation, clinical reasoning
FI-W12
Objective Renal/Infection/Hemodynamic Trend Summary
HD 1-6
Lab trends: Cr, K+, WBC, BP, I&Os
FI-W13
Medication Administration Record (MAR)
HD 1-6
What actually happened (not what should happen)
FI-W14
Nephrology Consultation
HD 3-4
Renal perspective, AKI management, medication recommendations
FI-W15
Cardiology Consultation
HD 3-4
Cardiac perspective, HF management, medication recommendations
FI-W16
Endocrinology Consultation
HD 4
Steroid physiology interpretation, adrenal risk (not proof of insufficiency)
FI-W17
Bedside Nursing Observation Notes
HD 1-6
Buried functional evidence, strength assessments, cognitive observations
FI-W18
Physical Therapy Assessment

HD 4-5
Strength, mobility, ambulation potential, functional trajectory
FI-W19
Occupational Therapy Assessment
HD 4-5
ADL capability, cognitive status during tasks, discharge readiness
FI-W20
Family Communication/Care Conference
HD 5
Mara Merrow (wife) concerns, family perspective, discharge wishes
FI-W21
Case Management/Social Work Discharge Planning
HD 5-6
Support systems, logistics, home equipment, follow-up coordination
FI-W22
Discharge-Facing Plan Snapshot
HD 5-7 (<05/23 18:00)
Visible-but-incomplete: organized improvement + deferred items
Properties preserved in all 22 files: - Closed-chart documentation (snapshots, not living plans) -
Temporal boundary: No file post-dates 05/23 18:00 - Multi-document synthesis required for all traps -
Source-hierarchy conflicts embedded (not resolved) - Evidence complete for all task demands

2.3 Task- Context Files (FI-T01-T07)

7 files total, all clinician-voice requests, no answers

File
Title
Task Workflow
Key Files
Anchor
FI-T01
Discharge Medication Reconciliation Request
Medication Reconciliation
FI-W01-W06, W13, W20
HD 5-7
FI-T02
Discharge Summary Drafting Request
Discharge Summary
Full ecosystem
HD 5-7
FI-T03
Discharge Readiness/Care Coordination Request
Discharge Planning
FI-W17-W21

HD 5-7
FI-T04
Consultant Synthesis/Interdisciplinary Plan Request
Care Coordination
FI-W14-W16, W20
HD 5-7
FI-T05
Early Post-Discharge Follow-Up Assessment
Post-Discharge Monitoring
Snapshot only
+7 days (no invented facts)
FI-T06
Patient Safety Readmission Risk Review
Safety Review
Snapshot only
+30 days (retrospective, not RCA)
FI-T07
Medication Safety Handoff Addendum
Medication Reconciliation support
Folded into FI-T01
HD 5-7
Properties: - Physician-authored clinician-voice framing - No answers or expected outputs embedded -
Named FI-W dependencies - Force synthesis across multiple files - Task-level traps (if any) distinct
from world-level

2.4 Supplementary Files (FI-S01-S04)

4 files total, all background/texture, non-load-bearing

File
Title
Content Type
Criticality
FI-S01
Remote PCI/Coronary Stent Provenance Summary
Background
CAD history only
FI-S02
Remote Sleep Study/OSA Provenance Summary
Background
Explicitly NOT for AMS explanation
FI-S03
Home Support/Equipment Logistics Reference
Logistics
Discharge environment context
FI-S04
Problem List/Past History Snapshot
Reference
Stale provenance, low-authority

Explicit boundary: Background, provenance, or logistics texture only. Must not carry sole critical evidence, resolve a friction, complete FI-W22, or become an answer source.

Ratio check: 4 supplementary / 33 total = ~12% supplementary (safe range)

2.5 World Spec AutoQC v6.3

113-check validation framework

Evidence: Applied to locked world specification

Result: All checks passed or documented with physician approval

Phase 3: Task Layer Construction

3.1 Task Prompts (TP-KM01-06)

6 prompts total, all physician-authored, clinician-voice

TP-KM01: Discharge Medication Reconciliation Request

Clinician-facing request to synthesize medication source conflicts, reconcile contradictions, propose discharge medication plan with source awareness and monitoring strategy

TP-KM02: Discharge Summary Drafting Request

Clinician-facing request to reconstruct patient's hospital course, document resolution and persistent issues, synthesize clinical reasoning

TP-KM03: Discharge Readiness/Care Coordination Assessment

Clinician-facing request to assess readiness for discharge, evaluate physiologic improvement vs. transition capability, coordinate care team

TP-KM04: Consultant Synthesis/Interdisciplinary Care Plan Request

Clinician-facing request to synthesize consultant perspectives without hierarchy collapse, develop unified care plan

TP-KM05: Early Post-Discharge Follow-Up Assessment (+7 days)

Prospective surveillance request, anchored at +7 post-discharge, no invented facts allowed

TP-KM06: Patient Safety Readmission Risk Review (+30 days)

Retrospective safety analysis, anchored at +30 post-discharge, explicitly NOT a root-cause analysis

3.2 Expected Outputs (EO-KM01-06)

6 expected outputs total, physician-authored, define reasoning bands

Properties: - Not goldens themselves; define strong reasoning without writing answers - Anti-overanswer (don't prescribe too narrowly) - Anti-underanswer (don't allow vague conclusions) - Two-sided friction protection - Name domains to explore, not answers to reach

3.3 Golden Responses (Golden-KM01-06)

6 goldens total, physician-authored, strong reference answers

Golden-KM01 (Medication Reconciliation)

"Leaves room for more than one defensible final medication sequence" key requirement is intentional, monitored, source-aware, communicated"

Golden-KM02 (Discharge Summary)

Reconstructs course HD1â†'HD3â†'HD4â†'HD5-6, "improved-but-unresolved" frame, multiple contributors

Golden-KM03 (Discharge Readiness)

"Discharge readiness is conditional" (not authorization); family AND team both defensible

Golden-KM04 (Interdisciplinary Care Plan)

Reconciles Cardiology and Nephrology without choosing winner; "steroid contribution" does not by itself prove adrenal insufficiency"

Golden-KM05 (+7 Follow-Up)

Surveillance frame, “should not invent post-discharge symptoms, labs, services, or outcomes”

Golden-KM06 (+30 Safety Review)

“Prospective readmission-risk synthesis,” explicitly “not a root-cause analysis”

Properties preserved in all six: - Uncertainty and multi-path defensibility maintained - No single-sequence logic (medication or course) - Both sides of all frictions remain two-sided - All hierarchies applied as reasoning, not shortcuts - FI-W22 visible-but-incomplete acknowledged - No adrenal insufficiency presented as revealed answer

3.4 Grader Guidance (GG-KM01-06)

6 grader-guidance files total, physician-authored, enforce scoring standards

File structure (all six files): 1. **Golden Reference:** States golden is benchmark, not key 2. **Good**

Practice That May Receive Credit But Should Not Be Required: Explicitly allows defensible

alternatives 3. **Known Errors:** Lists verbatim-matching and other failures to penalize 4. **Construction**

Notes: Confirms no points/thresholds/bands assigned

Anti-Verbatim-Matching Enforcement

Level 1 - Golden Reference (all six files):

Each states golden is benchmark, not a key. Example: “The golden is not a rigid medication answer key. Credit clinically safe, source-aware, task-appropriate medication reasoning even when wording, order, or exact sequencing differs.”

Level 2 - Good Practice (all six files):

Each ends with explicit anti-rigidity. Example: “Do not require one exact ordering of medication categories if the answer preserves the core safety logic.”

Level 3 - Known Errors (all six files):

Each names verbatim-matching itself as an error. Example: “Treating the golden as a fixed medication sequence that must be copied verbatim.”

Level 4 - Construction Notes (all six files):

Each reaffirms no scoring logic: “This document assigns no points, thresholds, or pass/fail bands.”

Multi-Path Defensibility Protections

GG-KM01 (Medication): - Credits “deliberate sequencing and monitoring” (not a single consultant winner) - Penalizes both “restarting all at once” AND “leaving all held stopped indefinitely” (preserves middle band) - Preserves rheumatology as highest prednisone authority without auto-deferring

GG-KM02 (Discharge Summary): - Credits different wording/ordering if course faithfully summarized - Penalizes “copy-forward sepsis-only framing” - Penalizes “steroids explain everything”

GG-KM03 (Discharge Readiness): - Distinguishes physiologic improvement from transition readiness - Credits conditional conclusions (not binary safe/unsafe) - Protects family evidence as meaningful, not automatically dispositive - Protects team perspective as also defensible

GG-KM04 (Interdisciplinary Care Plan): - Preserves cardiology vs nephrology tension - Penalizes “selecting Cardiology or Nephrology as automatically correct” - Rewards synthesis through trends/patient-status/timing - Keeps consultant disagreement resolution to synthesis, not rank

GG-KM05 (+7 Follow-Up): - Prohibits invented +7 facts (symptoms, labs, visits, outcomes) - Preserves transition-risk surveillance frame - Anchors to snapshot date; no post-close leakage

GG-KM06 (+30 Safety Review): - Prohibits invented outcomes/readmissions - Prohibits RCA drift - Prohibits blame assignment - Narrowed (per audit Finding-C) to endocrine as background context only

Friction Protection (All Three)

GG-KM01/04:

Penalize selecting consultant “as automatically correct” / “automatic winner”

GG-KM03:

Penalize ignoring nursing, PT, OT, CM/SW, family evidence; family concern “meaningful but not automatically dispositive”

GG-KM06:

Explicitly penalize “turning Endocrinology vs Primary Team into the dominant story rather than background steroid-source/endocrine-risk context”

Phase 4: Validation & Lock

4.1 Golden Architecture Validation

Date: May 31, 2026

Scope: Independent review of Golden-KM01-06 against locked brainstorm, world spec, and governance

Reviewer: Independent validation of physician-authored content

Validated properties: - Trap preservation (all 5 latent; no labels exposed) - Friction preservation (all 3 two-sided; no collapse) - Hierarchy preservation (all 3 hierarchies as reasoning; no shortcuts) - Adrenal insufficiency constraint (“meaningful but not proven” maintained across all six) - Source-hierarchy conflicts embedded (not resolved) - FI-W22 visible-but-incomplete consistency - FI-S non-load-bearing consistency - Multi-path defensibility maintained in all reasoning domains

Finding A - Validation Overclaim (Corrected):

Validation initially overclaimed source basis for some verification claims. Corrected by adding transparent annotation noting broader claim basis. Documentation-integrity correction, no clinical/architectural defect.

Finding B - Stale Continuity Surfaces (Corrected):

Governance surfaces stated “CANDIDATE REVIEW” while ratified as “LOCKED.” Updated to align with ratification record; metadata correction, no governance change.

Finding C - Golden-KM06 Endocrinology Friction Scope (Corrected):

Endocrinology-vs-Team friction was initially too broad in Golden-KM06. Corrected to “background steroid-source / endocrine-risk context only” to align with FI-T06’s permitted scope. Reflected across GG-KM06.

Lock Decision: RATIFIED AND LOCKED

4.2 Grader Guidance Construction Validation

Date: June 2, 2026

Scope: Independent review of GG-KM01-06 against locked architecture, goldens, expected outputs, prompts, and governance

11 validation objectives verified:

Objective 1: Count and Mapping (RESOLVED)

Six GG-KM files (KM01-06), no GG-KM07, one-to-one mapping with TP-KM01-06, EO-KM01-06, and Golden-KM01-06. FI-T07 folded into GG-KM01 addendum. Full TP → EO → Golden → GG chain verified.

Objective 2: Anti-Verbatim-Matching (RESOLVED)

Enforced at four independent levels per file: - Golden Reference: golden is benchmark, not key - Good Practice: explicitly allows defensible alternatives - Known Errors: lists verbatim-matching as penalizable error - Construction Notes: no points/thresholds/bands
Future graders can reasonably credit clinically sound alternatives.

Objective 3: Multi-Path Defensibility (RESOLVED)

Alternative defensible reasoning remains creditable; no hidden answer-key behavior; no single-sequence medication/course logic; conditional conclusions credited; both sides all three frictions remain defensible.

Objective 4: Medication Guidance (GG-KM01) (RESOLVED)

Cardiology vs Nephrology tension preserved; prednisone hierarchy preserved (rheumatology strongest); inpatient-vs-outpatient distinction preserved; staged reasoning rewarded; does NOT become medication answer key; does NOT force one restart sequence; does NOT over-credit consultant-winner selection.

Objective 5: Discharge Readiness (GG-KM03) (RESOLVED)

Physiologic improvement distinguished from transition readiness; family concern appropriately handled (meaningful evidence, not automatically dispositive); PT/OT/nursing/CM/SW evidence protected. Does NOT become discharge-authorization rubric; does NOT create hidden disposition key; does NOT force one conclusion.

Objective 6: Follow-Up and Outcome Protection (RESOLVED)

GG-KM05: prohibits invented +7 facts, preserves transition-risk framing

GG-KM06: prohibits invented outcomes/readmissions, prohibits RCA drift, prohibits blame, honors anchor-as-review-date rule

Objective 7: Hierarchy Protection (RESOLVED)

All three hierarchies preserved as reasoning, not shortcuts. Prednisone hierarchy with rheumatology highest protected. Source hierarchy explicitly NOT a shortcut for clinical disagreement resolution.

Objective 8: Friction Protection (RESOLVED)

All three frictions preserved with anti-winner guardrails. GG-KM06 correctly scopes endocrine as background per audit Finding-C. Guidance rewards defensible synthesis, penalizes winner-selection.

Objective 9: Source Fidelity (RESOLVED)

FI-T stays task framing, FI-W stays evidence, FI-S stays background/support. FI-W22 stays visible-but-incomplete across all six. No FI-S becomes sole evidence. No score leakage.

Objective 10: Grading Leakage Boundary (RESOLVED)

Guidance does NOT become scoring rubric, point system, threshold system, pass/fail framework. Every GG file's Construction Notes explicitly assigns no points/thresholds/bands. Scoring reserved for separate, later-authorized phase.

Objective 11: Validation Review Accuracy (RESOLVED)

Validation accurately reflects GG files; one honest disclosure: architecture-footer status corrected from CANDIDATE REVIEW to LOCKED (metadata alignment, not substantive change). Documented transparently.

True Defects: NONE

Governance/Metadata Inconsistencies: NONE residual (Finding-C corrected and reflected)

Physician Preference Items: 1. Good Practice sections are key multi-path protectors; ensure these stay genuinely optional in eventual scoring 2. GG-KM01 Known Errors list appropriately strict on prednisone-source and inpatient-lispro (weight as safety-critical without rigid key)

Future AutoQC Watch Items: 1. Official Section 6 Grader Guidelines AutoQC prompt not yet imported (must be imported before grader-guidance AutoQC runs) 2. Scoring construction (points/thresholds/pass-fail) reserved for separate, later-authorized phase 3. Keep GG-KM03/05/06 scoring distinct (present readiness / prospective follow-up / retrospective safety) 4. Packaging-layer World Spec AutoQC carry-forwards remain live (2.52 ratio, 2.113 Tool/Origin, 2.108 administrative-deliverable, v6.3 8-column file plan)

Lock Decision: LOCK READY - All six GG-KM files approved for ratification and lock

Guidance Quality: STRONG

Benchmark Integrity: HIGH

Final Recommendation: GO

4.3 Audit Findings & Reconciliations

Finding A - Validation Overclaim (Documentation-Integrity Correction) - Issue: Validation review initially overclaimed source basis for verification claims - **Resolution:** Transparent annotation added noting broader claim basis - **Type:** Documentation-integrity correction, not clinical/architectural defect - **Impact:** Audit disclosure maintains transparency; no hidden corrections

Finding B - Stale Continuity Surfaces (Metadata Alignment) - Issue: Current-facing governance surfaces stated "CANDIDATE REVIEW" while actual governance ratified to "LOCKED" - **Resolution:** Updated surface status to match ratification record - **Type:** Metadata-alignment correction, no substantive governance change - **Impact:** Documentation now consistent with locked decisions

Finding C - Golden-KM06 Endocrinology Friction Scope (Architecture Refinement) - Issue: Endocrinology-vs-Team friction initially too broad in Golden-KM06 - **Resolution:** Narrowed to "background steroid-source / endocrine-risk context only" to align with FI-T06 permitted scope - **Impact:** Reflected across all GG-KM06 content; grader guidance explicitly penalizes over-amplification

Finding D - FI-S Boundary Clarification - Issue: FI-S (supplementary) files under-specified in initial guidance - **Resolution:** Explicit wording: "background, provenance, or logistics texture only" must not carry sole critical evidence, resolve a friction, complete FI-W22, or become an answer source" - **Impact:** 4 supplementary files now clearly non-load-bearing; 29 essential files protected as primary evidence

Reconciliation Discipline: All findings disclosed transparently rather than made silently. Audit documents explain findings, resolutions, and impact on downstream layers. This transparent approach establishes documentation-integrity discipline across all future phases.

4.4 Lock-Readiness Confirmation

Date: June 2, 2026

Lock Decision: LOCK READY

Complete Ecosystem Locked and Verified

22 world-level files (FI-W01-W22) - LOCKED

7 task-context files (FI-T01-T07) - LOCKED

4 supplementary files (FI-S01-S04) - LOCKED

6 task prompts (TP-KM01-06) - LOCKED

6 expected outputs (EO-KM01-06) - LOCKED

6 golden responses (Golden-KM01-06) - LOCKED

6 grader-guidance documents (GG-KM01-06) - LOCKED

Governance and validation documentation - LOCKED

All Key Properties Preserved

Trap integrity: All 5 traps remain latent, reasoning-required, no labels exposed

Friction integrity: All 3 frictions remain two-sided with anti-winner guardrails

Hierarchy integrity: All 3 hierarchies preserved as reasoning, not shortcuts

Synthesis requirement: Every task forces multi-document reconciliation

FI-W22: Visible-but-incomplete (incomplete without FI-W series, FI-T context, consultant details)

FI-S: Non-load-bearing (background/provenance/logistics only, never sole evidence)

Source-of-truth conflicts: Embedded, not resolved

Multi-path defensibility: Alternative defensible approaches credited at all layers

Anti-verbatim-matching: Enforced four times per GG file

AutoQC Readiness

Section 2 World Spec AutoQC v6.3 (113 checks): Ready to run

Section 6 Grader Guidelines AutoQC: Prompt must be imported before running

Packaging Carry-Forwards (Standing)

2.52 Essential/Supplementary Ratio: Verify 29 essential / 4 supplementary $\hat{=}$ 12% supplementary

2.113 Tool/Origin Two-Track Specification: Assign per-file engineering-converted vs writer-produced-media origin with filename convention

2.108 Administrative Deliverable: Confirm Discharge Planning Documentation satisfies requirement (utilization-review named fallback)

v6.3 8-Column File Plan: Ensure Section 3 uses v6.3 structure (Source + Tool as separate columns)

Forward Dependencies

AutoQC Phase (Authorization Required)

Blocking Issue

Section 6 Grader Guidelines AutoQC prompt not yet in local repository. Must be imported from official Anthropic reference before grader-guidance AutoQC can run.

AutoQC Tasks (In Sequence)

Import Section 6 Grader Guidelines AutoQC prompt

Run Section 2 World Spec AutoQC v6.3 (113 checks) against locked world spec

Run Section 6 Grader Guidelines AutoQC against locked grader guidance (GG-KM01-06)

Confirm all checks pass or document physician approvals

Packaging Phase (Dependent on AutoQC Completion)

Carry-Forward 1: 2.52 Essential/Supplementary Ratio

Current: 29 essential (22 world + 7 task-context), 4 supplementary $\hat{=}$ 12%

Requirement: Confirm against 2.52 specification (safe direction)

Action: Verify ratio at packaging

Carry-Forward 2: 2.113 Tool/Origin Two-Track Filename Convention

Current: Per-file Tool/Origin specification needed

Requirement: Assign engineering-converted vs writer-produced-media origin per 2.113

Action: Apply two-track naming convention at packaging

Carry-Forward 3: 2.108 Administrative Deliverable

Current: Discharge Planning Documentation (TP-KM03/T05/T06 under 4-workflow architecture)

Requirement: Confirm administrative-deliverable requirement satisfied

Fallback: Utilization-review named if reviewer challenges

Action: Finalize at packaging

Carry-Forward 4: v6.3 8-Column File Plan (Section 3)

Current: File plan needs v6.3 structure (Source, Tool as separate columns)

Requirement: Rebuild Section 3 table with v6.3 format (not 05/06 template's 7 columns)

Action: Rebuild Section 3 at packaging with v6.3 structure

END OF TRANSCRIPT