

Raising Task Difficulty: Worked Example #1

Task Details

A skilled nursing facility admission note for a complicated long-stay ICU patient discharged from acute care.

The original world files contained a discharge summary from the acute stay that leaked the traps. With the world files left unchanged and no task-level traps, **the model got all the required non-negotiables and content correct.**

Disclaimer: Scores are just for reference, the outputs before are after showed clear model failure over trajectories when outputs were reviewed.

Before	97% 1	95% 2	95% 3	98% 4	97% 5	97% 6	95% 7	96% 8	98% 9	96% 10
After	35% 1	25% 2	40% 3	92% 4	30% 5	15% 6	25% 7	30% 8	32% 9	78% 10

Change 1: Take the answer key out of the world files

The original world files included the hospital discharge summary. It summarized the entire course and effectively laid out the traps and their answers, so the model only had to transcribe it. The discharge summary was removed from the world files. If a later task genuinely needs that document, it can be added back as a task-level file scoped to that task rather than as a shared world file. **Removing this file forced the model to reconcile information across the stack of other patient documents.** To close the loop, the prompt stated that the discharge summary had not yet been faxed to the SNF provider, but that the patient's chart was available to refer to.

Change 2: Add noise and require a format

A SNF admission template was added as a task-level file, with an instruction that the response follow it. A nursing intake note containing some information that did not contribute to the task's non-negotiables was also added as noise. **The model then has to fit its reasoning into a required structure and sieve out what actually matters from the nursing intake.**

Change 3: Bury a critical finding off the text

A task-level photograph of the patient's PICC line, clearly infected, was added. The only other hints were a low-grade fever in the nursing intake vitals and a nursing note recording that the photo had been taken, without any description of it.

This change did most of the work, for four reasons:

- Images are fair game for the model to read, but it does not read them well.
- The new issue appears nowhere in prose, only in the photo and one low-grade temperature, which mirrors how these things surface in practice.
- It is a genuinely critical finding. Missing an infected PICC in a patient receiving IV vancomycin through that same line would be indefensible in real care.
- The grader was tightened to match: missing it is a hard error and an immediate fail, regardless of how strong the rest of the note is.



Task-level image created by Nanobanana Pro.

Quick Summary

To take any clinical documentation task from easy to hard.

1 Pull the answer key. Remove or relocate any source that states the conclusion. World files should provide raw material, not the synthesis.

2 Force reconciliation. No single file should be complete, so the model must cross-check sources to get it right.

3 Don't spell out the trap. As in real practice, leave it off the text: an image, a single vital, or a buried line in a long note. Real signal is rarely highlighted.

4 Match the grader to the stakes. If a miss would be clinically unacceptable, make it a hard error and penalize it appropriately.

5 Keep it true to practice. The traps that hold up are the ones that mirror how information actually goes missing.

Other tips

- Images help a lot (made with Nanobanana Pro): medication bottles, handwritten scribbles of patient med lists, resident notes during rounds, clinic notes during downtime, or an infected line as in this case.

- Audio files, even though they need to be accompanied by transcripts, seem to give nice failures.

- Preliminary imaging reports that suggest something like a filling defect in a particular artery, which a physician would pick up automatically but the model often does not.

- Bundling labs (clearly dated) into a single task-level or world-level file instead of providing them separately.

- Guideline timing: if the world is set in 2024 with a 2024 date in the prompt, any guideline used should be from that time. If a guideline changed afterward (say 2025), check whether the model still cites the 2024 version or reaches for the 2025 one, which fails timelessness.

- Adding several task-level files (more than five), not all of them relevant to the task, which adds noise.

- Asking the model to fill in a form or template.

- Psych-related tasks seem to work well, mainly because there is less objective information (like labs) for the task to depend on.

Niranjana M

12 June 2026